

### **Multinomial Naive Bayes:**

Training on enron1 dataset

Results for enron1

Accuracy: 0.9320

Precision: 0.9275

Recall: 0.8591

F1-score: 0.8920

Training on enron2 dataset

Results for enron2

Accuracy: 0.9351

Precision: 0.9091

Recall: 0.8462

F1-score: 0.8765

Training on enron4 dataset

Results for enron4

Accuracy: 0.9687

Precision: 0.9698

Recall: 0.9872

F1-score: 0.9785

### **Discrete Naive Bayes:**

Test Set Performance:

Accuracy: 0.9013

Precision: 0.9906

Recall: 0.7047

F1-score: 0.8235

### **Logistic Regression:**

Processing dataset: enron1

Training with  $\lambda = 0.01$

Iteration 0: Log-Likelihood = -243.1827

Iteration 100: Log-Likelihood = -4.1844

Iteration 200: Log-Likelihood = -2.5338

Iteration 300: Log-Likelihood = -1.9531

Iteration 400: Log-Likelihood = -1.6655

F1-score for  $\lambda=0.01$ : 0.8941

Training with  $\lambda = 0.1$

Iteration 0: Log-Likelihood = -244.9118

Iteration 100: Log-Likelihood = -7.4283

Iteration 200: Log-Likelihood = -6.2140

Iteration 300: Log-Likelihood = -5.8135

Iteration 400: Log-Likelihood = -5.6027

F1-score for  $\lambda=0.1$ : 0.8941

Training with  $\lambda = 1$

Iteration 0: Log-Likelihood = -262.2027

Iteration 100: Log-Likelihood = -22.1016

Iteration 200: Log-Likelihood = -20.6634

Iteration 300: Log-Likelihood = -20.5250

Iteration 400: Log-Likelihood = -20.5185

F1-score for  $\lambda=1$ : 0.8571

Training with  $\lambda = 10$

Iteration 0: Log-Likelihood = -435.1118

Iteration 100: Log-Likelihood = -60.6955

Iteration 200: Log-Likelihood = -60.6955

Iteration 300: Log-Likelihood = -60.6955

Iteration 400: Log-Likelihood = -60.6955

F1-score for  $\lambda=10$ : 0.8889

Best  $\lambda$ : 0.01 with F1-score: 0.8941

Iteration 0: Log-Likelihood = -478.6928

Iteration 100: Log-Likelihood = -5.4579

Iteration 200: Log-Likelihood = -3.3377

Iteration 300: Log-Likelihood = -2.5907

Iteration 400: Log-Likelihood = -2.2203

Results for enron1 (bow representation):

Accuracy: 0.9583

Precision: 0.9062

Recall: 0.9732

F1-score: 0.9385

Training with  $\lambda = 0.01$

Iteration 0: Log-Likelihood = -123.2335

Iteration 100: Log-Likelihood = -5.8988

Iteration 200: Log-Likelihood = -3.4078

Iteration 300: Log-Likelihood = -2.5564

Iteration 400: Log-Likelihood = -2.1417

F1-score for  $\lambda=0.01$ : 0.9630

Training with  $\lambda = 0.1$

Iteration 0: Log-Likelihood = -123.4203

Iteration 100: Log-Likelihood = -8.8273

Iteration 200: Log-Likelihood = -7.2598

Iteration 300: Log-Likelihood = -6.9001

Iteration 400: Log-Likelihood = -6.7813

F1-score for  $\lambda=0.1$ : 0.9630

Training with  $\lambda = 1$

Iteration 0: Log-Likelihood = -125.2889  
Iteration 100: Log-Likelihood = -27.2217  
Iteration 200: Log-Likelihood = -27.1360  
Iteration 300: Log-Likelihood = -27.1321  
Iteration 400: Log-Likelihood = -27.1318  
F1-score for  $\lambda=1$ : 0.9630

Training with  $\lambda = 10$

Iteration 0: Log-Likelihood = -143.9748  
Iteration 100: Log-Likelihood = -78.3121  
Iteration 200: Log-Likelihood = -78.3121  
Iteration 300: Log-Likelihood = -78.3121  
Iteration 400: Log-Likelihood = -78.3121  
F1-score for  $\lambda=10$ : 0.9091

Best  $\lambda$ : 0.01 with F1-score: 0.9630

Iteration 0: Log-Likelihood = -208.2585  
Iteration 100: Log-Likelihood = -7.0405  
Iteration 200: Log-Likelihood = -4.0790  
Iteration 300: Log-Likelihood = -3.0730  
Iteration 400: Log-Likelihood = -2.5844

Results for enron1 (bernoulli representation):

Accuracy: 0.9583  
Precision: 0.9221  
Recall: 0.9530  
F1-score: 0.9373

Processing dataset: enron2

Training with  $\lambda = 0.01$

Iteration 0: Log-Likelihood = -167.5701  
Iteration 100: Log-Likelihood = -6.2887  
Iteration 200: Log-Likelihood = -3.9561  
Iteration 300: Log-Likelihood = -3.0781  
Iteration 400: Log-Likelihood = -2.6249  
F1-score for  $\lambda=0.01$ : 0.9538

Training with  $\lambda = 0.1$

Iteration 0: Log-Likelihood = -169.3095  
Iteration 100: Log-Likelihood = -9.7804  
Iteration 200: Log-Likelihood = -8.0364  
Iteration 300: Log-Likelihood = -7.4520  
Iteration 400: Log-Likelihood = -7.1601  
F1-score for  $\lambda=0.1$ : 0.9394

Training with  $\lambda = 1$

Iteration 0: Log-Likelihood = -186.7039  
Iteration 100: Log-Likelihood = -25.5403  
Iteration 200: Log-Likelihood = -23.8745  
Iteration 300: Log-Likelihood = -23.6900  
Iteration 400: Log-Likelihood = -23.6752  
F1-score for  $\lambda=1$ : 0.9231

Training with  $\lambda = 10$

Iteration 0: Log-Likelihood = -360.6479  
Iteration 100: Log-Likelihood = -65.9558  
Iteration 200: Log-Likelihood = -65.9558  
Iteration 300: Log-Likelihood = -65.9558  
Iteration 400: Log-Likelihood = -65.9558  
F1-score for  $\lambda=10$ : 0.8667

Best  $\lambda$ : 0.01 with F1-score: 0.9538

Iteration 0: Log-Likelihood = -396.6470  
Iteration 100: Log-Likelihood = -7.7300  
Iteration 200: Log-Likelihood = -4.9763  
Iteration 300: Log-Likelihood = -3.9291  
Iteration 400: Log-Likelihood = -3.3855

Results for enron2 (bow representation):

Accuracy: 0.9519  
Precision: 0.8963  
Recall: 0.9308  
F1-score: 0.9132

Training with  $\lambda = 0.01$

Iteration 0: Log-Likelihood = -120.7644  
Iteration 100: Log-Likelihood = -6.7985  
Iteration 200: Log-Likelihood = -4.2010  
Iteration 300: Log-Likelihood = -3.2582  
Iteration 400: Log-Likelihood = -2.7766  
F1-score for  $\lambda=0.01$ : 0.9333

Training with  $\lambda = 0.1$

Iteration 0: Log-Likelihood = -120.9552  
Iteration 100: Log-Likelihood = -9.6300  
Iteration 200: Log-Likelihood = -7.9616  
Iteration 300: Log-Likelihood = -7.5379  
Iteration 400: Log-Likelihood = -7.3840  
F1-score for  $\lambda=0.1$ : 0.9189

Training with  $\lambda = 1$

Iteration 0: Log-Likelihood = -122.8636  
Iteration 100: Log-Likelihood = -27.4583  
Iteration 200: Log-Likelihood = -27.3644

Iteration 300: Log-Likelihood = -27.3603  
Iteration 400: Log-Likelihood = -27.3600  
F1-score for  $\lambda=1$ : 0.9167

Training with  $\lambda = 10$   
Iteration 0: Log-Likelihood = -141.9474  
Iteration 100: Log-Likelihood = -77.9907  
Iteration 200: Log-Likelihood = -77.9907  
Iteration 300: Log-Likelihood = -77.9907  
Iteration 400: Log-Likelihood = -77.9907  
F1-score for  $\lambda=10$ : 0.8657

Best  $\lambda$ : 0.01 with F1-score: 0.9333  
Iteration 0: Log-Likelihood = -212.9513  
Iteration 100: Log-Likelihood = -9.1320  
Iteration 200: Log-Likelihood = -5.6136  
Iteration 300: Log-Likelihood = -4.3237  
Iteration 400: Log-Likelihood = -3.6672

Results for enron2 (bernoulli representation):  
Accuracy: 0.9498  
Precision: 0.9077  
Recall: 0.9077  
F1-score: 0.9077

Processing dataset: enron4

Training with  $\lambda = 0.01$   
Iteration 0: Log-Likelihood = -225.9457  
Iteration 100: Log-Likelihood = -3.5305  
Iteration 200: Log-Likelihood = -2.0956  
Iteration 300: Log-Likelihood = -1.6194  
Iteration 400: Log-Likelihood = -1.3903  
F1-score for  $\lambda=0.01$ : 0.9798

Training with  $\lambda = 0.1$   
Iteration 0: Log-Likelihood = -227.0083  
Iteration 100: Log-Likelihood = -6.5138  
Iteration 200: Log-Likelihood = -5.5225  
Iteration 300: Log-Likelihood = -5.2305  
Iteration 400: Log-Likelihood = -5.0839  
F1-score for  $\lambda=0.1$ : 0.9798

Training with  $\lambda = 1$   
Iteration 0: Log-Likelihood = -237.6343  
Iteration 100: Log-Likelihood = -21.7123  
Iteration 200: Log-Likelihood = -20.7580  
Iteration 300: Log-Likelihood = -20.6659  
Iteration 400: Log-Likelihood = -20.6578

F1-score for  $\lambda=1$ : 0.9758

Training with  $\lambda = 10$   
Iteration 0: Log-Likelihood = -343.8945  
Iteration 100: Log-Likelihood = -67.5273  
Iteration 200: Log-Likelihood = -67.5273  
Iteration 300: Log-Likelihood = -67.5273  
Iteration 400: Log-Likelihood = -67.5273  
F1-score for  $\lambda=10$ : 0.9680

Best  $\lambda$ : 0.01 with F1-score: 0.9798  
Iteration 0: Log-Likelihood = -513.1457  
Iteration 100: Log-Likelihood = -4.1940  
Iteration 200: Log-Likelihood = -2.5208  
Iteration 300: Log-Likelihood = -1.9624  
Iteration 400: Log-Likelihood = -1.6926

Results for enron4 (bow representation):  
Accuracy: 0.9595  
Precision: 0.9467  
Recall: 1.0000  
F1-score: 0.9726

Training with  $\lambda = 0.01$   
Iteration 0: Log-Likelihood = -169.0172  
Iteration 100: Log-Likelihood = -4.3566  
Iteration 200: Log-Likelihood = -2.5046  
Iteration 300: Log-Likelihood = -1.9049  
Iteration 400: Log-Likelihood = -1.6199  
F1-score for  $\lambda=0.01$ : 0.9808

Training with  $\lambda = 0.1$   
Iteration 0: Log-Likelihood = -169.3462  
Iteration 100: Log-Likelihood = -7.2094  
Iteration 200: Log-Likelihood = -6.0944  
Iteration 300: Log-Likelihood = -5.8622  
Iteration 400: Log-Likelihood = -5.7882  
F1-score for  $\lambda=0.1$ : 0.9808

Training with  $\lambda = 1$   
Iteration 0: Log-Likelihood = -172.6365  
Iteration 100: Log-Likelihood = -25.1857  
Iteration 200: Log-Likelihood = -25.1145  
Iteration 300: Log-Likelihood = -25.1121  
Iteration 400: Log-Likelihood = -25.1120  
F1-score for  $\lambda=1$ : 0.9808

Training with  $\lambda = 10$   
Iteration 0: Log-Likelihood = -205.5387

Iteration 100: Log-Likelihood = -80.1764  
Iteration 200: Log-Likelihood = -80.1764  
Iteration 300: Log-Likelihood = -80.1764  
Iteration 400: Log-Likelihood = -80.1764  
F1-score for  $\lambda=10$ : 0.9808

Best  $\lambda$ : 0.01 with F1-score: 0.9808  
Iteration 0: Log-Likelihood = -285.3286  
Iteration 100: Log-Likelihood = -5.5128  
Iteration 200: Log-Likelihood = -3.1834

Iteration 300: Log-Likelihood = -2.4233  
Iteration 400: Log-Likelihood = -2.0613

Results for enron4 (bernoulli representation):  
Accuracy: 0.9705  
Precision: 0.9607  
Recall: 1.0000  
F1-score: 0.9799

Each classifier's performance was optimized through hyperparameter tuning; for Naive Bayes classifiers, the main hyperparameter, alpha, regulates Laplace smoothing; for Logistic Regression, the key hyperparameter is the regularization strength (C), which was varied across a logarithmic scale from 0.01 to 10; the optimization solver used was 'lbfgs'; and the maximum number of iterations was set to 1000 to ensure convergence.

1. Which combination of algorithm and data representation yielded the best performance? Why?

The combination of the Bag of Words representation with Logistic Regression performed the best. This method produced the best accuracy and F1-score, most likely as a result of the model's capacity to identify subtle word associations in the dataset. Logistic regression can model more intricate interactions, improving predictive performance, whereas Naive Bayes assumes that features are independent.

2. Did Multinomial Naive Bayes perform better than Logistic Regression on the Bag of Words representation? Explain.

No, using the Bag of Words representation, Multinomial Naive Bayes did not perform better than Logistic Regression. Multinomial Naive Bayes assumes that word frequencies are conditionally independent given the class, which makes it ideal for text classification tasks. In contrast, logistic regression is more effective for the given dataset because it does not rely on this assumption and is capable of learning more intricate decision boundaries.

3. Did Discrete Naive Bayes perform better than Logistic Regression on the Bernoulli representation? Explain.

No, Discrete Naive Bayes (Bernoulli Naive Bayes) did not outperform Logistic Regression on the Bernoulli representation. Although Bernoulli Naive Bayes is designed for binary feature representations, it struggles with correlated features, which are common in natural language processing tasks. Logistic Regression is more robust in handling such dependencies and thus demonstrated superior performance in terms of accuracy and F1-score.